

Hive:

```
[cloudera@quickstart ~]$ hive
Logging initialized using configuration in jar:file:/usr/jars/hive-common-1.1.0-
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> show databases;
OK
default
demol
emp
Time taken: 0.638 seconds, Fetched: 3 row(s)
```

```
hive> drop database emp;
OK
Time taken: 0.112 seconds
hive> create database emp;
OK
Time taken: 0.077 seconds
hive> create table emp.employees(EmpId int, Age int, sales_percent float, No_of_pro
ject int, increment float)
> row format delimited
> fields terminated by ',' ;
OK
Time taken: 0.273 seconds
hive> describe emp.employees;
OK
empid                int
age                  int
sales_percent        float
no_of_project        int
increment            float
Time taken: 0.112 seconds, Fetched: 5 row(s)
```

```
hive> create table if not exists emp.copy_table like emp.employees;
OK
Time taken: 0.125 seconds
```

```
hive> load data local inpath '/home/cloudera/Downloads/sales_data_sample.csv' in
to table emp.employees;
Loading data to table emp.employees
Table emp.employees stats: [numFiles=1, totalSize=371]
```

OK

Time taken: 0.57 seconds

```
hive> select * from emp.employees;
```

OK

10107	30	95.7	2	2871.0
10121	34	81.35	5	2765.9
10134	41	94.74	2	3884.34
10145	45	83.26	6	3746.7
10159	49	100.0	14	5205.27
10168	36	96.66	1	3479.76
10180	29	86.13	9	2497.77
10188	48	100.0	1	5512.32
10201	22	98.57	2	2168.54
10211	41	100.0	14	4708.44
10223	37	100.0	1	3965.66
10237	23	100.0	7	2333.12
10251	28	100.0	2	3188.64
10263	34	100.0	2	3676.76

```
hive> select * from emp.employees2;
```

OK

10107	30	95.7	2	2871.0
10121	34	81.35	5	2765.9
10134	41	94.74	2	3884.34
10145	45	83.26	6	3746.7
10159	49	100.0	14	5205.27
10168	36	96.66	1	3479.76
10180	29	86.13	9	2497.77
10188	48	100.0	1	5512.32
10201	22	98.57	2	2168.54
10211	41	100.0	14	4708.44
10223	37	100.0	1	3965.66
10237	23	100.0	7	2333.12
10251	28	100.0	2	3188.64
10263	34	100.0	2	3676.76
10275	45	92.83	1	4177.35

Time taken: 0.088 seconds, Fetched: 15 row(s)

```
hive> Alter table emp.employees rename to emp.employees1;
```

OK

Time taken: 0.116 seconds

```
hive> use emp;
OK
Time taken: 0.011 seconds
hive> show tables;
OK
copy_table
employees1
employees2
Time taken: 0.011 seconds, Fetched: 3 row(s)
```

```
hive> select employees1.empid,employees1.empid from employees1,employees2;
Warning: Map Join MAPJOIN[7][bigTable=employees2] in task 'Stage-3:MAPRED' is a cross product
Query ID = cloudera_20221108021515_c2be0502-19ff-4712-b602-cb37da41bc32
Total jobs = 1
Execution log at: /tmp/cloudera/cloudera_20221108021515_c2be0502-19ff-4712-b602-cb37da41bc32.log
2022-11-08 02:15:47 Starting to launch local task to process map join; maximum memory = 49807360
2022-11-08 02:15:48 Dump the side-table for tag: 0 with group count: 1 into file: file:/tmp/cloudera/35025124-0251-42d3-9bce-da8388c5b69a/hive_2022-11-08_02-15-43_644_3264878058206245337-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable
2022-11-08 02:15:48 Uploaded 1 File to: file:/tmp/cloudera/35025124-0251-42d3-9bce-da8388c5b69a/hive_2022-11-08_02-15-43_644_3264878058206245337-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable (392 bytes)
2022-11-08 02:15:48 End of local task; Time Taken: 1.124 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1667898921734_0001, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1667898921734_0001/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1667898921734_0001
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2022-11-08 02:15:58,541 Stage-3 map = 0%, reduce = 0%
2022-11-08 02:16:04,856 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.82 sec
MapReduce Total cumulative CPU time: 1 seconds 820 msec
```

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 1.82 sec HDFS Read: 6434 HDFS Write: 2700 SUCCESS

Total MapReduce CPU Time Spent: 1 seconds 820 msec

OK

10107 10107

10121 10121

10134 10134

10145 10145

10159 10159

10168 10168

10180 10180

10188 10188

10201 10201

10211 10211

10223 10223

10237 10237

10251 10251

10263 10263

10275 10275

10107 10107

10121 10121

10134 10134

10145 10145

10251	10251
10263	10263
10275	10275
10107	10107
10121	10121
10134	10134
10145	10145
10159	10159
10168	10168
10180	10180
10188	10188
10201	10201
10211	10211
10223	10223
10237	10237
10251	10251
10263	10263
10275	10275

Time taken: 22.333 seconds, Fetched: 225 row(s)

Apache pig :

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
Job DAG:  
job_local1501966295_0002  
  
2022-11-14 23:27:49,502 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!  
2022-11-14 23:27:49,503 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS  
2022-11-14 23:27:49,503 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum  
2022-11-14 23:27:49,504 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been initialized  
2022-11-14 23:27:49,551 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1  
2022-11-14 23:27:49,551 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1  
(1,utkarsh,gupta,,pune)  
(2,parth,kabadi,,mumbai)  
(3,harsh,pande,,amravati)  
(4,aryan,bhusari,,pune)  
grunt> describe MyCSVData  
MyCSVData: {id: int,firstname: chararray,lastname: chararray,phone: int,city: chararray}  
grunt> █
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
Output(s):  
Successfully stored records in: "file:/tmp/temp-1306075918/tmp-1499326309"  
  
Job DAG:  
job_local1501966295_0002  
  
2022-11-14 23:27:49,502 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!  
2022-11-14 23:27:49,503 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS  
2022-11-14 23:27:49,503 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum  
2022-11-14 23:27:49,504 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been initialized  
2022-11-14 23:27:49,551 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1  
2022-11-14 23:27:49,551 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1  
(1,utkarsh,gupta,,pune)  
(2,parth,kabadi,,mumbai)  
(3,harsh,pande,,amravati)  
(4,aryan,bhusari,,pune)  
grunt> █
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
ation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum  
2022-11-14 23:17:12,973 [main] INFO org.apache.hadoop.conf.Configuration.deprec  
ation - fs.default.name is deprecated. Instead, use fs.defaultFS  
2022-11-14 23:17:12,981 [main] INFO org.apache.hadoop.conf.Configuration.deprec  
ation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum  
2022-11-14 23:17:13,059 [main] INFO org.apache.hadoop.conf.Configuration.deprec  
ation - fs.default.name is deprecated. Instead, use fs.defaultFS  
2022-11-14 23:17:13,061 [main] INFO org.apache.hadoop.conf.Configuration.deprec  
ation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum  
2022-11-14 23:17:13,183 [main] INFO org.apache.hadoop.conf.Configuration.deprec  
ation - fs.default.name is deprecated. Instead, use fs.defaultFS  
2022-11-14 23:17:13,185 [main] INFO org.apache.hadoop.conf.Configuration.deprec  
ation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum  
2022-11-14 23:17:13,281 [main] INFO org.apache.hadoop.conf.Configuration.deprec  
ation - fs.default.name is deprecated. Instead, use fs.defaultFS  
2022-11-14 23:17:13,282 [main] INFO org.apache.hadoop.conf.Configuration.deprec  
ation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum  
grunt> MyCSVData = LOAD '/home/cloudera/Downloads/sample_1.csv' using PigStorage  
(', ') AS (id:int,firstname:chararray,lastname:chararray,phone:int,city:chararray  
);  
grunt> █
```

```
*sample_1-2.csv (~/.Downloads) - gedit  
File Edit View Search Tools Documents Help  
Open Save Undo  
*sample_1-2.csv  
1,utkarsh,gupta,9860762595,pune  
2,parth,kabadi,8778903465,mumbai  
3,harsh,pande,7889057481,amravati  
4,aryan,bhusari,9809813345,pune  
Plain Text Tab Width: 8 Ln 4, Col 32 INS
```



```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System Mon Nov 14, 11:31 PM cloudera
cloudera@quickstart:~
File Edit View Search Terminal Help
onconcatenation threshold: 100 optimistic? false
2022-11-14 23:31:14,583 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer
- MR plan size before optimization: 1
2022-11-14 23:31:14,583 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer
- MR plan size after optimization: 1
2022-11-14 23:31:14,584 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the
job
2022-11-14 23:31:14,584 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler
- mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2022-11-14 23:31:14,596 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been ini
tialized
2022-11-14 23:31:14,599 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigMapOnly$Map - Al
ias being processed per job phase (AliasName[line,offset]): M: MyCSVData[2,12] C: R:
2022-11-14 23:31:14,599 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File c
onconcatenation threshold: 100 optimistic? false
2022-11-14 23:31:14,600 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer
- MR plan size before optimization: 1
2022-11-14 23:31:14,600 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer
- MR plan size after optimization: 1
2022-11-14 23:31:14,601 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the
job
2022-11-14 23:31:14,601 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler
- mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2022-11-14 23:31:14,603 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been ini
tialized
2022-11-14 23:31:14,666 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigMapOnly$Map - Al
ias being processed per job phase (AliasName[line,offset]): M: MyCSVData[2,12] C: R:
2022-11-14 23:31:14,666 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File c
onconcatenation threshold: 100 optimistic? false
2022-11-14 23:31:14,667 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer
- MR plan size before optimization: 1
2022-11-14 23:31:14,667 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer
- MR plan size after optimization: 1
2022-11-14 23:31:14,667 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the
job
2022-11-14 23:31:14,667 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler
- mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2022-11-14 23:31:14,683 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been ini
tialized
2022-11-14 23:31:14,685 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigMapOnly$Map - Al
ias being processed per job phase (AliasName[line,offset]): M: MyCSVData[2,12] C: R:
-----
| MyCSVData | id:int | firstname:chararray | lastname:chararray | phone:int | city:chararray |
-----
|           | 4      | aryan               | bhusari            | 9809813345 | pune           |
-----

grunt>
```



```
cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
Applications Places System Mon Nov 14, 11:30 PM cloudera@quickstart:~
cloudera@quickstart:~
File Edit View Search Terminal Help

| |
| |---Project[bytearray][4] - scope-51
| |
| |---MyCSVData: Load(/home/cloudera/Downloads/sample_1-2.csv:PigStorage(','))
- scope-38

2022-11-14 23:29:36,125 [main] INFO org.apache.pig.backend.hadoop.executionengi
ne.mapReduceLayer.MRCCompiler - File concatenation threshold: 100 optimistic? fal
se
2022-11-14 23:29:36,128 [main] INFO org.apache.pig.backend.hadoop.executionengi
ne.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2022-11-14 23:29:36,128 [main] INFO org.apache.pig.backend.hadoop.executionengi
ne.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
#-----
# Map Reduce Plan
#-----
MapReduce node scope-56
Map Plan
MyCSVData: Store(fakefile:org.apache.pig.builtin.PigStorage) - scope-55
|
|---MyCSVData: New For Each(false,false,false,false,false)[bag] - scope-54
| |
| |Cast[int] - scope-40
| | |
| | |---Project[bytearray][0] - scope-39
| | |
| |Cast[chararray] - scope-43
| | |
| | |---Project[bytearray][1] - scope-42
| | |
| |Cast[chararray] - scope-46
| | |
| | |---Project[bytearray][2] - scope-45
| | |
| |Cast[int] - scope-49
| | |
| | |---Project[bytearray][3] - scope-48
| | |
| |Cast[chararray] - scope-52
| | |
| | |---Project[bytearray][4] - scope-51
| |
| |---MyCSVData: Load(/home/cloudera/Downloads/sample_1-2.csv:PigStorage(','))
- scope-38-----
Global sort: false
-----
grunt> █
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
Output(s):  
Successfully stored records in: "file:/tmp/temp1223820747/tmp-1156799919"  
  
Job DAG:  
job_local1016883113_0003  
  
2022-11-14 23:35:17,997 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!  
2022-11-14 23:35:17,998 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS  
2022-11-14 23:35:17,998 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - io.bytes.per.checksum is deprecated. Instead, use dfs.bytes-per-checksum  
2022-11-14 23:35:17,998 [main] WARN org.apache.pig.data.SchemaTupleBackend - SchemaTupleBackend has already been initialized  
2022-11-14 23:35:18,004 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1  
2022-11-14 23:35:18,004 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1  
(pune,{(4,aryan,bhusari,,pune),(1,utkarsh,gupta,,pune)})  
(mumbai,{(2,parth,kabadi,,mumbai)})  
(amravati,{(3,harsh,pande,,amravati)})  
grunt> █
```